

Bayesian Semiparametric Regression Based on MCMC Simulation Techniques: A Tutorial

Thomas Kneib, Stefan Lang and Andreas Brezger

Department of Statistics, University of Munich.

February 10, 2009

Abstract

This tutorial demonstrates the usage of *BayesX* for analysing Bayesian semiparametric regression models based on MCMC simulation techniques. As an example, we consider data on undernutrition of children in Zambia. The tutorial is designed to be self-contained and describes all features of *BayesX* in detail, that will be needed throughout the tutorial. Therefore it may also serve as a first introduction into the general usage of *BayesX*. Note that the tutorial is written for the version of *BayesX* including a graphical user interface and graphics facilities. While all regression-related commands can be executed in the command line version as well, the graphics facilities will not be available.

Contents

1	Introduction	2
2	Description of the data set	2
3	Getting started	3
4	Reading data set information	3
5	Map objects	5
6	Bayesian semiparametric regression	9
7	Visualising estimation results	15
7.1	Post estimation commands	15
7.2	Graph Objects	16
8	Customising graphics	17
9	Autocorrelation functions and sampling paths	20
10	Sensitivity analysis	21
	References	22

1 Introduction

This tutorial demonstrates the usage of *BayesX* for analysing Bayesian semiparametric regression models based on MCMC simulation techniques. As an example, we consider data on childhood undernutrition in Zambia. This data has already been analysed in Kandala et al. (2001) and we will use the same model that has been developed there. Since our focus is on demonstrating how regression models can be estimated in *BayesX*, we do not discuss or interpret the estimation results but simply give the commands to obtain them.

The main focus in this tutorial is on full Bayesian inference based on MCMC techniques. *BayesX* also supports two further inferential concepts: empirical Bayes inference based on mixed model methodology and a method for simultaneously selecting variables and smoothing parameters, which are described in two additional tutorials. All tutorials are designed to be self-contained and describe all features of *BayesX* in detail, that will be needed throughout the tutorial. Users who are already familiar with the usage of *dataset* and *map objects* may therefore skim through sections 3–5.

The theoretical background of Bayesian semiparametric regression will not be described in this tutorial. The reference manual may serve as an introduction, full details about the estimation techniques for the full Bayesian approach can be found in Fahrmeir & Lang (2001a), Fahrmeir & Lang (2001b), Lang & Brezger (2004) and Brezger & Lang (2006). Survival models are treated in Hennerfeind, Brezger & Fahrmeir (2006) while count data regression is covered in Fahrmeir & Osuna (2006).

2 Description of the data set

Undernutrition among children is usually determined by assessing an anthropometric status of the children relative to a reference standard. In our example, undernutrition is measured by stunting or insufficient height for age, indicating chronic undernutrition. Stunting for a child i is determined using a Z-score defined as

$$Z_i = \frac{AI_i - MAI}{\sigma}$$

where AI refers to the child’s anthropometric indicator (height at a certain age in our example), while MAI and σ correspond to the median and the standard deviation in the reference population, respectively.

Our main interest is on modelling the dependence of undernutrition on covariates including the age of the child, the body mass index of the child’s mother, the district the child lives in and some further categorical covariates. Table 1 gives a description of the variables that we will use in our model.

Variable	Description
<i>hazstd</i>	standardised Z-score for stunting
<i>bmi</i>	body mass index of the mother
<i>age</i>	age of the child in months
<i>district</i>	district where the mother lives
<i>rcw</i>	mother’s employment status with categories “working” (= 1) and “not working” (= -1)
<i>edu1/2</i>	mother’s educational status with categories “complete primary but incomplete secondary” (<i>edu1</i> = 1), “complete secondary or higher” (<i>edu2</i> = 1) and “no education or incomplete primary” (<i>edu1</i> = <i>edu2</i> = -1)
<i>tpr</i>	locality of the domicile with categories “urban” (= 1) and “rural” (= -1)
<i>sex</i>	gender of the child with categories “male” (= 1) and “female” (= -1)

Table 1: Variables in the undernutrition data set.

3 Getting started

After having started the graphical user interface version of *BayesX*, a main window with four sub-windows appears on the screen. These are the *command window* for entering and executing code, the *output window* for displaying results, the *review window* for easy access to past commands, and the *object browser* that displays all objects currently available. In the command line version of *BayesX*, there are, of course, no sub-windows but only command line prompt to enter commands.

BayesX is object oriented although the concept is limited, i.e. inheritance and other concepts of object oriented languages like C++ or R are not supported. For every object type, a number of object-specific methods can be applied to a particular object instance. The syntax for generating a new object in *BayesX* is

```
> objecttype objectname
```

where *objecttype* defines the type of the object to be created, e.g. **dataset**, and *objectname* is the name to be assigned to the new object.

The rest of the tutorial is separated in seven parts dealing with the different steps of estimating a regression model with MCMC simulation techniques. In section 4, we create a *dataset object* to store, handle and manipulate the data. We will also give a brief description of some methods that may be applied to *dataset objects*. Since we want to estimate a spatial effect of the district in which a child lives, we need the boundaries of the districts to compute the neighbourhood information of the map of Zambia. This information will be stored in a *map object*. Section 5 describes how to create and handle these objects. Estimation of the regression model is carried out in section 6 using a *bayesreg object*. The next two sections describe how to visualise the estimation results and how to customise the obtained graphics. Section 9 describes post estimation commands which can be used to investigate the sampling paths and the autocorrelation functions of the estimated parameters. In a last section we perform a sensitivity analysis to assess the impact of hyperparameter choices on our estimation results.

If you have not done so yet, please download the data set and the *boundary file* associated with this tutorial now (from the *BayesX* homepage). You may also want to download the batch file containing the commands used in the following sections. Please note, that paths within these commands must be changed according to the storage location of the corresponding files on your hard disk.

4 Reading data set information

In a first step, we read the available data set information into *BayesX*. Therefore we create a *dataset object* named **d**:

```
> dataset d
```

We store the data in **d** using the method **infile**:

```
> d.infile, maxobs=5000 using c:\data\zambia.raw
```

Note, that we assume the data to be provided in the external file **c:\data\zambia.raw**. The first few lines of this file look like this:

```
hazstd bmi agc district rcw edu1 edu2 tpr sex
0.0791769 21.83 4 81 -1 1 0 1 -1
-0.2541965 21.83 26 81 -1 1 0 1 -1
-0.1599823 20.43 56 81 1 -1 -1 1 1
0.1733911 22.27 6 81 -1 0 1 1 1
```

In our example, the file contains the variable names in the first line. Therefore, it is not necessary to specify them in the `infile` command. If the file contained only the data without variable names, we would have to supply them after the keyword `infile`:

```
> d.infile hazstd bmi agc district rcw edu1 edu2 tpr sex, maxobs=5000
    using c:\data\zambia.raw
```

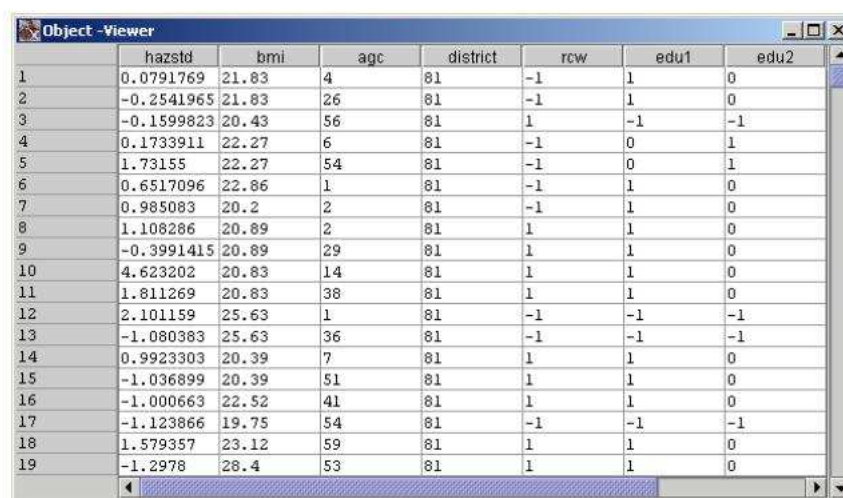
Option `maxobs` can be used to speed up the execution time of the `infile` command. If `maxobs` is specified, *BayesX* allocates enough memory to store all the data while the total amount of required memory is unknown in advance if `maxobs` remains unspecified. For larger data sets, this may cause *BayesX* to start reading the data set information several times because the currently allocated memory is exceeded. However, this is only meaningful for larger data sets with more than 10,000 observations and could therefore be omitted in our example.

A second option that may be added to the `infile` command is the `missing` option to indicate missing values. Specifying for example `missing = M` defines the letter 'M' as an indicator for a missing value. The default for missing values are a period '.' and 'NA' (which remain valid indicators for missing values even if an additional indicator is defined by the `missing` option).

After having read the dataset information, we can inspect the data visually. Executing the command

```
> d.describe
```

opens an *Object-Viewer* window containing the data in form of a spreadsheet (see Figure 1). The same can also be achieved by double-clicking on the *dataset object* in the *object browser*.



	hazstd	bmi	agc	district	rcw	edu1	edu2
1	0.0791769	21.83	4	81	-1	1	0
2	-0.2541965	21.83	26	81	-1	1	0
3	-0.1599823	20.43	56	81	1	-1	-1
4	0.1733911	22.27	6	81	-1	0	1
5	1.73155	22.27	54	81	-1	0	1
6	0.6517096	22.86	1	81	-1	1	0
7	0.985083	20.2	2	81	-1	1	0
8	1.108286	20.89	2	81	1	1	0
9	-0.3991415	20.89	29	81	1	1	0
10	4.623202	20.83	14	81	1	1	0
11	1.811269	20.83	38	81	1	1	0
12	2.101159	25.63	1	81	-1	-1	-1
13	-1.080383	25.63	36	81	-1	-1	-1
14	0.9923303	20.39	7	81	1	1	0
15	-1.036899	20.39	51	81	1	1	0
16	-1.000663	22.52	41	81	1	1	0
17	-1.123866	19.75	54	81	-1	-1	-1
18	1.579357	23.12	59	81	1	1	0
19	-1.2978	28.4	53	81	1	1	0

Figure 1: A screenshot of the dataset.

Further methods allow to examine the variables in the *dataset object*. For a categorical variable such as *sex*, the `tabulate` command may be used to produce a frequency table:

```
> d.tabulate sex
```

resulting in

Variable: sex

Value	Obs	Freq	Cum
-1	2451	0.5057	0.5057

1	2396	0.4943	1
---	------	--------	---

being printed in the *output window*. For continuous variables, the `descriptive` command prints several characteristics of the variable in the output window. E.g., executing

```
> d.descriptive bmi
```

leads to

Variable	Obs	Mean	Median	Std	Min	Max
bmi	4847	21.944349	21.4	3.2879659	12.8	39.29

5 Map objects

In the following, we will estimate a spatially correlated effect of the district in which a child lives. Therefore we need the boundaries of the districts in Zambia to compute the neighbourhood information of the map of Zambia. We therefore create a *map object*

```
> map m
```

and read the boundaries using the `infile` command of *map objects*:

```
> m.infile using c:\data\zambia.bnd
```

Having read the boundary information, *BayesX* automatically computes the neighbourhood matrix of the map.

The file following the keyword `using` is assumed to contain the boundaries in form of closed polygons. To give an example we print a small part of the boundary file of Zambia. The map corresponding to the section of the boundary file can be found in Figure 2.

```

:
"52",48
28.080507,-12.537530
28.083376,-12.546980
28.109501,-12.548961
28.134972,-12.566787
28.154797,-12.585320
28.165771,-12.593912
28.165771,-12.593912
28.160769,-12.609917
28.152800,-12.633824
28.144831,-12.657733
28.132877,-12.677656
28.120922,-12.701565
28.120922,-12.717505
28.120922,-12.741411
28.116938,-12.761335
28.108969,-12.777274
28.100998,-12.793213
28.089045,-12.817122
28.085060,-12.837045
28.081076,-12.856968
28.081076,-12.876892
28.080862,-12.884153
28.080862,-12.884153
28.076630,-12.879521
28.031454,-12.881046
27.974281,-12.884675
27.910725,-12.878692
27.686228,-12.880120
27.665676,-12.854732
27.653563,-12.818301

```

```

27.639263,-12.759848
27.648254,-12.699927
27.662464,-12.680613
27.662464,-12.680613
27.666534,-12.675080
27.703260,-12.679779
27.752020,-12.695455
27.797932,-12.702188
27.836775,-12.707567
27.867813,-12.699892
27.902308,-12.667418
27.922668,-12.630853
27.943035,-12.596350
27.963434,-12.571486
27.983179,-12.563844
28.016331,-12.554779
28.070650,-12.542199
28.080507,-12.537530

```

```

:
:

```

For each region of the map, the boundary file must contain the identifying name of the region, the polygons that form the boundary of the region, and the number of lines the polygon consists of. The first line always contains the region code surrounded by quotation marks and the number of lines the polygon of the region consists of. The code and the number of lines must be separated by a comma. The subsequent lines contain the coordinates of the straight lines that form the boundary of the region. The straight lines are represented by the coordinates of their end points. Coordinates must be separated by a comma. Note that the first and the last point must be identical (see the example above) to obtain a closed polygon. Compare chapter 5 of the reference manual for a detailed description of some special cases, e.g. regions divided into subregions.




Figure 2: Corresponding graph of the section of the boundary file

Map objects may be visualised using method `describe`:

```
> m.describe
```

resulting in the graph shown in Figure 3. Additionally, `describe` prints further information about the *map object* in the *output window* including the name of the object, the number of regions, the minimum and maximum number of neighbours and the bandwidth of the corresponding adjacency or neighbourhood matrix:

```

MAP m
Number of regions: 54
Minimum number of neighbors: 1
Maximum number of neighbors: 9
Bandsize of corresponding adjacency matrix: 24

```

The numerical complexity associated with the estimation of structured spatial effects using MCMC




Figure 3: The districts within Zambia.

techniques depends essentially on the structure of the neighbourhood matrix. Often the geographical information stored in a boundary file does not represent the “ideal” ordering of the districts or regions (with respect to the estimation problem). Therefore it may be useful to reorder the map using method `reorder`:

```
> m.reorder
```

Usually, reordering results in a smaller bandwidth although the bandwidth is not the criterion that is minimised by `reorder`. Instead the *envelope* of the neighbourhood matrix is minimised (compare George & Liu (1981)).

In order to avoid reordering the *map object* every time you start *BayesX*, it is useful to store the reordered version in a separate file. This can be achieved using the `outfile` command of *map objects*:

```
> m.outfile, replace using c:\data\zambiasort.bnd
```

The reordered map is now stored in the given file. Note, that specifying the option `replace` allows *BayesX* to overwrite an existing file with the same name. Without this option, an error message would be raised if the given file is already existing.

Reading the boundary information from an external file and computing the neighbourhood matrix may be a computationally intensive task if the map contains a large number of regions or if the polygons are given in great detail. To avoid doing these computation in every *BayesX* session, we store the neighbourhood information in a *graph file* using method `outfile` together with the `graph` option:

```
> m.outfile, replace graph using c:\data\zambiasort.gra
```

A graph file stores the nodes and the edges of a graph $G = (N, E)$, see for example George & Liu (1981) for a first introduction into graph theory. A graph is a convenient way of representing the neighbourhood structure of a geographical map. The nodes of the graph correspond to the region codes. The neighbourhood structure is represented by the edges of the graph. In some situations it may be useful to define weights associated with the edges of a graph which can be stored in the *graph file* as well.

We now describe the structure of a graph file as it is expected by *BayesX*. The first line of a *graph file* must contain the total number of nodes of the graph. In the remaining lines, the nodes of the graph together with their edges and associated weights are specified. One node corresponds to three consecutive lines. The first of the three lines must contain the name of the node, which typically will be the name of the geographical region. In the second line, the number of edges of that particular node is given. The third line contains the corresponding edges of the node, where

an edge is given by the index of a neighbouring node. The index starts with zero. For example, if the fourth and the seventh node/region in the *graph file* are connected/neighbours, the edge index for the fourth node/region is 6 and for the seventh node/region 3.

We illustrate the structure of a graph file with an example. The following few lines are the beginning of the graph file corresponding to the reordered map of Zambia:

```
57
87
1
5
76
3
9 8 7
67
2
10 9
```

⋮

The first line specifies the total number of nodes, in the present example 57 nodes. The subsequent three lines correspond to the node with name ‘87’, which is the first region in the reordered map of Zambia. Region ‘87’ has 1 neighbour, namely the sixth node appearing in the graph file. Once again, note that the index starts with zero, i.e. 0 corresponds to the first node, 1 corresponds to the second node and so on. Lines 5 to 7 in the example correspond to node ‘76’ and its three neighbours and lines 8 to 10 correspond to node ‘67’.

In a graph file it is also possible to specify weights associated with the edges of the nodes. Since in the preceding example no weights are explicitly specified, all weights are automatically defined to be equal to one. Nonequal weights are specified in the graph file by simply adding them following the edges of a particular node. An example of the beginning of a graph file with weights is given below:

```
57
87
1
5 1.44172
76
3
7 8 9 0.707424 1.3816 0.682372
67
2
9 10 1.67424 0.8406
```

⋮

Here the edge of the first node ‘87’ has weight 1.44172, the edges of the second node have weights 0.707424, 1.3816 and 0.682372.

Note, that graph files allow the estimation of very general correlated effects based on Markov random fields. While the polygons stored in a *boundary file* represent geographical information, the nodes and edges of a graph may define arbitrary neighbourhood structures. For example, the definition of three-dimensional Markov random fields representing space-time interactions is possible.

To see how storing maps in *graph files* affects the computation time of the `infile` command, we create a second *map object* and read in the information from the graph file. Again, we have to specify the keyword `graph`:


```
> map m1
> m1.infile, graph using c:\data\zambiasort.gra
```

As you should have noticed, reading geographical information from a *graph file* is usually much faster than reading from a *boundary file*. However, using *graph files* also has a drawback. Since they do no longer contain the full information on the polygons forming the map, we can not visualise a *map object* created from a *graph file*. Trying to do so

```
> m1.describe
```

raises an error message. This implies, that visualising estimation results of spatial effects can only be based on *map objects* created from *boundary files*, although estimation can be carried out using *graph files*. Since we will work with the *map object* `m` in the following, we delete `m1`:

```
> drop m1
```

6 Bayesian semiparametric regression

To estimate a regression model using MCMC simulation techniques, we first create a *bayesreg object*:

```
> bayesreg b
```

By default, estimation results are written to the subdirectory **output** of the installation directory. In this case, the default filenames are composed of the name of the *bayesreg object* and the type of the specific file. Usually it is more convenient to store the results in a user-specified directory. To define this directory, we use the **outfile** command of *bayesreg objects*:

```
> b.outfile = c:\data\b
```

Note, that **outfile** does not only specify a directory but also a base filename (the character 'b' in our example). Therefore, executing the command above leads to storage of the results in the directory `c:\data` and all filenames start with the character 'b'. Of course, the base filename may be different from the name of the *bayesreg object*.

In addition to parameter estimates, *BayesX* also gives acceptance rates for the different effects and some further information on the estimation process. In contrast to parameter estimates, this information is not stored automatically but is printed in the *output window*. Therefore it is useful to store the contents of the *output window*. This can be achieved automatically by opening a *log file* using the **logopen** command

```
> logopen, replace using c:\data\logmcmc.txt
```

After opening a *log file*, every information written to the output window is also stored in this file. Option **replace** allows *BayesX* to overwrite an existing file with the same name as the specified *log file*. Without **replace** results are appended to an existing file.

The model presented in Kandala et al. (2001) is given by the following semiparametric predictor:

$$\eta = \gamma_0 + \gamma_1 rcw + \gamma_2 edu1 + \gamma_3 edu2 + \gamma_4 tpr + \gamma_5 sex + f_1(bmi) + f_2(agg) + f^{str}(district) + f^{unstr}(district).$$

The two continuous covariates *bmi* and *agg* are assumed to have a possibly nonlinear effect on the Z-score and are therefore modelled nonparametrically (as P-splines with second order random walk prior in our example). The spatial effect of the district is split up into a spatially correlated part $f^{str}(district)$ and an uncorrelated part $f^{unstr}(district)$, see Fahrmeir & Lang (2001b) for a motivation. The correlated part is modelled by a Markov random field prior, where the neighbourhood matrix and possible weights associated with the neighbours are obtained from the *map object* `m`. The uncorrelated part is modelled by an i.i.d. Gaussian effect.

To estimate the model we use method **regress** of *bayesreg objects*:

```
> b.regress hazstd = rcw + edu1 + edu2 + tpr + sex + bmi(psplinerw2)
+ agc(psplinerw2) + district(spatial,map=m) + district(random),
family=gaussian iterations=12000 burnin=2000 step=10 predict using d
```

Options `iterations`, `burnin` and `step` define properties of the MCMC algorithm. The total number of MCMC iterations is given by `iterations` while the number of burn in iterations is given by `burnin`. Therefore we obtain a sample of 10000 random numbers with the above specifications. Since, in general, these random numbers are correlated, we do not use all of them but thin out the Markov chain by the thinning parameter `step`. Specifying `step=10` as above forces *BayesX* to store only every 10th sampled parameter which leads to a random sample of length 1000 for every parameter in our example. Note, that the choice of `iterations` of course also affects computation time.

If option `predict` is specified, samples of the deviance, the effective number of parameters p_D , and the deviance information criteria *DIC* of the model are computed, see Spiegelhalter et al. (2002). In addition, estimates for the linear predictor and the expectation of every observation are obtained.

In the following, we reproduce the content of the *output window* to make the user familiar with the estimation results produced by *BayesX*. Note that the output may look somewhat different depending on the version of *BayesX* you are considering.

ESTIMATION RESULTS:

Predicted values:

Estimated mean of predictors, expectation of response and individual deviances are stored in file
c:\data\b_predictmean.raw

Estimation results for the deviance:

Unstandardized Deviance (-2*Loglikelihood(y|mu))

Mean:	12688.726
Std. Dev:	12.993345
2.5% Quantile:	12663.772
10% Quantile:	12672.535
50% Quantile:	12688.173
90% Quantile:	12705.219
97.5% Quantile:	12714.784

Saturated Deviance (-2*Loglikelihood(y|mu) + 2*Loglikelihood(y|mu=y))

Mean:	4850.7212
Std. Dev:	98.805525
2.5% Quantile:	4658.3601
10% Quantile:	4719.064
50% Quantile:	4854.0429
90% Quantile:	4981.005
97.5% Quantile:	5046.1822

Samples of the deviance are stored in file
c:\data\b_deviance_sample.raw

Estimation results for the DIC:

DIC based on the unstandardized deviance

Deviance(bar_mu):	12639.017
pD:	49.708864
DIC:	12738.435

DIC based on the saturated deviance

Deviance(bar_mu): 4800.0024
pD: 50.718841
DIC: 4901.44

Estimation results for the scale parameter:

Acceptance rate: 100 %

Mean: 0.80205
Std. dev.: 0.016386
2.5% Quantile: 0.7711
10% Quantile: 0.780591
50% Quantile: 0.801434
90% Quantile: 0.824332
97.5% Quantile: 0.835445

FixedEffects1

Acceptance rate: 100 %

Variable	mean	Std. Dev.	2.5% quant.	median	97.5% quant.
const	0.102144	0.0468836	0.0115916	0.101498	0.196307
rcw	0.00780289	0.0136821	-0.0186219	0.00754262	0.0340403
edu1	-0.0622753	0.0265637	-0.116141	-0.0613026	-0.0125551
edu2	0.235591	0.0465318	0.142865	0.234172	0.325396
tpr	0.0898013	0.0216879	0.0483689	0.0901556	0.132811
sex	-0.0585167	0.0125313	-0.0826825	-0.0583656	-0.0339112

Results for fixed effects are also stored in file
c:\data\b_FixedEffects1.res

f_bmi_pspline

Acceptance rate: 100 %

Results are stored in file
c:\data\b_f_bmi_pspline.res

Postscript file is stored in file
c:\data\b_f_bmi_pspline.ps

Results may be visualized using method 'plotnonp'
Type for example: objectname.plotnonp 1

f_bmi_pspline_variance

Acceptance rate: 100 %

Estimation results for the variance component:

Mean: 0.00174445
Std. dev.: 0.0022271
2.5% Quantile: 0.000265649
10% Quantile: 0.000419949
50% Quantile: 0.00112203
90% Quantile: 0.00375514
97.5% Quantile: 0.00654175

Results for the variance component are also stored in file
c:\data\b_f_bmi_pspline_var.res

Estimation results for the smoothing parameter:

Mean:	937.489	(df: 3.35257)
Std. dev.:	773.149	
2.5% Quantile:	121.533	(df: 5.531)
10% Quantile:	211.716	(df: 4.85283)
50% Quantile:	717.565	(df: 3.58964)
90% Quantile:	1910.98	(df: 2.78449)
97.5% Quantile:	3038.87	(df: 2.46105)

Results for the smoothing parameter are also stored in file
c:\data\b_f_bmi_p spline_lambda.res

f_agc_p spline

Acceptance rate: 100 %

Results are stored in file
c:\data\b_f_agc_p spline.res

Postscript file is stored in file
c:\data\b_f_agc_p spline.ps

Results may be visualized using method 'plotnonp'
Type for example: objectname.plotnonp 3

f_agc_p spline_variance

Acceptance rate: 100 %

Estimation results for the variance component:

Mean:	0.00630648
Std. dev.:	0.00788173
2.5% Quantile:	0.00124029
10% Quantile:	0.00179477
50% Quantile:	0.0041396
90% Quantile:	0.0121621
97.5% Quantile:	0.0228385

Results for the variance component are also stored in file
c:\data\b_f_agc_p spline_var.res

Estimation results for the smoothing parameter:

Mean:	234.684	(df: 6.28883)
Std. dev.:	177.494	
2.5% Quantile:	34.9805	(df: 9.37168)
10% Quantile:	66.326	(df: 8.23766)
50% Quantile:	194.159	(df: 6.55631)
90% Quantile:	442.421	(df: 5.457)
97.5% Quantile:	647.435	(df: 5.0029)

Results for the smoothing parameter are also stored in file
c:\data\b_f_agc_p spline_lambda.res

f_district_spatial

Acceptance rate: 100 %

Results are stored in file

c:\data\b_f_district_spatial.res

Postscript file is stored in file
c:\data\b_f_district_spatial.ps

Results may be visualized in BayesX using method 'drawmap'
Type for example: objectname.drawmap 5

f_district_spatial_variance

Acceptance rate: 100 %

Estimation results for the variance component:

Mean:	0.0360104
Std. dev.:	0.0179351
2.5% Quantile:	0.0105888
10% Quantile:	0.0165585
50% Quantile:	0.0324881
90% Quantile:	0.0600347
97.5% Quantile:	0.0826229

Results for the variance component are also stored in file
c:\data\b_f_district_spatial_var.res

Estimation results for the smoothing parameter:

Mean:	28.7933	(df: 0)
Std. dev.:	16.8091	
2.5% Quantile:	9.73429	(df: 0)
10% Quantile:	13.3299	(df: 0)
50% Quantile:	24.7758	(df: 0)
90% Quantile:	48.715	(df: 0)
97.5% Quantile:	75.0528	(df: 0)

Results for the smoothing parameter are also stored in file
c:\data\b_f_district_spatial_lambda.res

f_district_random

Acceptance rate: 100 %

Results for random effects are stored in file
c:\data\b_f_district_random.res

Results for the sum of the structured and unstructured
spatial effects are stored in file
c:\data\b_district_spatialtotal.res

f_district_random_variance

Acceptance rate: 100 %

Estimation results for the variance component:

Mean:	0.00788548
Std. dev.:	0.00618319
2.5% Quantile:	0.000696144
10% Quantile:	0.00138013
50% Quantile:	0.00645178
90% Quantile:	0.0165836

97.5% Quantile: 0.0237649

Results for the variance component are also stored in file
c:\data\b_f_district_random_var.res

Files of model summary:

Batch file for visualizing effects of nonlinear functions is stored in file
c:\data\b_graphics.prg

NOTE: 'input filename' must be substituted by the filename of the boundary-file

Batch file for visualizing effects of nonlinear functions
in R / S-Plus is stored in file
c:\data\b_r_splus.txt

NOTE: 'input filename' must be substituted by the filename of the boundary-file

Latex file of model summaries is stored in file
c:\data\b_model_summary.tex

In addition to the information being printed to the *output window*, results for each effect are written to external ASCII files. The names of these files are given in the output window, compare the previous pages. The files contain the posterior mean and median, the posterior 2.5%, 10%, 90% and 97.5% quantiles, and the corresponding 95% and 80% posterior probabilities of the estimated effects. For example, the beginning of the file c:\data\b_f_bmi_pspline.res for the effect of *bmi* may look like this:

intnr	bmi	pmean	pqu2p5	pqu10	pmed	pqu90	pqu97p5	pcat95	pcat80
1	12.8	-0.284065	-0.660801	-0.51678	-0.283909	-0.0585753	0.085998	0	-1
2	13.15	-0.276772	-0.609989	-0.483848	-0.275156	-0.070517	0.0572406	0	-1
3	14.01	-0.258674	-0.515628	-0.416837	-0.257793	-0.10009	-0.00289024	-1	-1

The posterior quantiles and posterior probabilities may be changed by the user via the options `level1` and `level2`. For example, specifying `level1=99` and `level2=70` in the option list of the `regress` command leads to the computation of 0.5%, 15%, 85% and 99.5% quantiles of the posterior. The defaults are `level1=95` and `level2=80`.

Some nonparametric effects are visualised by *BayesX* automatically and the resulting graphs are stored in *ps* format. For example, the effect of *bmi* is visualised in the file c:\data\b_f_bmi_pspline.ps (compare the results on the previous pages for the other filenames). In addition to the *ps files* a file containing the commands to reproduce the graphics is stored in the output directory. In our example the name of the file is c:\data\b_graphics.prg. The advantage is that additional options may be added by the user to customise the graphs (compare the following two sections).

Moreover, a file with ending `.tex` is created in the outfile directory. This file contains a summary of the estimation results and may be compiled using \LaTeX .

Having finished the estimation, we may close the *log file* by typing

```
> logclose
```

Note, that the *log file* is closed automatically when you exit *BayesX*.

7 Visualising estimation results

BayesX provides three possibilities to visualise estimation results:

- As mentioned in the previous section, certain results are automatically visualised by *BayesX* and stored in *ps files*.
- Post estimation commands of *bayesreg objects* allow to visualise results after having executed a **regress** command.
- *Graph objects* may be used to produce graphics using the ASCII files containing the estimation results. In principle, *graph objects* allow the visualisation of any content of a *dataset object*. *Graph files* are also used in the batch file containing the commands to reproduce the automatically generated graphics.

In this section, we describe the general usage of the post estimation commands as well as the commands for the usage with *graph objects* to enable the user to reproduce the automatically generated plots directly in *BayesX*. Section 8 describes how to customise plots.

7.1 Post estimation commands

After having estimated a regression model, plots for nonparametric effects of metrical covariates can be produced using the post estimation command **plotnonp**:

```
> b.plotnonp 1
```

and

```
> b.plotnonp 3
```

produce the graphs shown in Figure 4 in an *object-viewer window*. The numbers following the **plotnonp** command depend on the order in which the model terms have been specified (and an internal ordering of the effect types). The numbers are supplied in the *output window* after estimation, compare the results in the previous section.

By default, the plots contain the posterior mean and pointwise credible intervals according to the levels specified in the **regress** command. So by default the plot includes pointwise 80% and 95% credible intervals.




Figure 4: Effect of the body mass index of the child's mother and of the age of the child together with pointwise 80% and 95% credible intervals.

A plot may be stored in *ps* format using the **outfile** option. Executing

```
> b.plotnonp 1, replace outfile = c:\data\f_bmi.ps
```

stores the plot for the estimated effect of *bmi* in the file `c:\data\f_bmi.ps`. Again, specifying `replace` allows *BayesX* to overwrite an existing file. Note that *BayesX* does not display the graph on the screen if the option `outfile` is specified.

Estimation results for spatial effects are best visualised by drawing the respective map and colouring the regions of the map according to some characteristic of the posterior, e.g. the posterior mean. For the structured spatial effect this can be achieved using the post estimation command `drawmap`

```
> b.drawmap 5
```

which results in the graph shown in Figure 5.




Figure 5: Posterior mean of the structured spatial effect.

7.2 Graph Objects

The commands presented in the previous subsection work only after having estimated a regression model in the current *BayesX* session. However, it may of course also be useful to visualise results of former analyses. This can be achieved using *graph objects*. Note again that *graph files* are also used in the batch file containing the commands to reproduce the automatically generated graphics. Therefore, the purpose of this subsection is also to enable the user to understand the content of this batch file.

In a first step, we read the estimation results into a *dataset object*. For example, the estimation results for the effect of *bmi* can be read into *BayesX* by executing the commands

```
> dataset res
> res.infile using c:\data\b_f_bmi_pspline.res
```

Now the estimation results (or any content of a *dataset object*) may be visualised using a *graph object* which we create by typing

```
> graph g
```

The results stored in the *dataset object* `res` are now visualised using the `plot` command of *graph objects*. Executing

```
> g.plot bmi pmean pqu2p5 pqu10 pqu90 pqu97p5 using res
```

reproduces the graph in Figure 4.

Similar as for `plotnonp`, the direct usage of the `drawmap` command is only possible after executing a `regress` command. However, using *graph objects* again allows us to visualise results that have been stored in a file.

First, we read the information contained in this file into a *dataset object*. For example, the following command

```
> res.infile using c:\data\b_f_district_spatial.res
```

stores the estimation results for the structured spatial effect in the *dataset object* `res`. Now we can visualise the posterior mean using method `drawmap` of *graph objects* leading again to the graph shown in Figure 5:

```
> g.drawmap pmean district, map=m using res
```

Since – in contrast to a *bayesreg object* – no *map object* is associated with a *graph object* we have to specify the map that we want to use explicitly in the option list.

Using *graph objects* also allows us to plot other characteristics of the posterior than the posterior mean. For instance the posterior 95% probabilities may be visualised by

```
> g.drawmap pcat95 district, map=m using res
```

The result is shown in Figure 6.




Figure 6: Posterior 95% probability of the structured spatial effect.

A further advantage of *graph objects* is, that they allow to visualise the estimation results for the uncorrelated spatial effects. Since these are modelled as unstructured random effects, *BayesX* is unable to recognise them as spatial effects. However, proceeding as follows gives us the possibility to plot the unstructured spatial effect shown in Figure 7:

```
> res.infile using c:\data\b_f_district_random.res
> g.drawmap pmean district, map=m using res
```

8 Customising graphics

This section describes how to customise graphics created in *BayesX*. All options are described for the usage with the post estimation commands but may be used with graph files as well. Hence, the options presented in this section also enable the user to modify the batch file containing the commands to reproduce the automatically generated graphics.




Figure 7: Posterior mean of the unstructured spatial effect.

For the presentation of nonparametric effects it may be desirable to include only one of the credible intervals into the plot. This is achieved by specifying the `levels` option. Possible values of this option are 1 and 2, corresponding to the levels specified in the `regress` command (compare section 6). If the default values of `level1` and `level2` have been used, specifying `level=2` in the `plotnonp` command causes *BayesX* to plot the 80% credible interval only (Figure 8):

```
> b.plotnonp 1, levels=2
```




Figure 8: Effect of the body mass index of the child's mother with pointwise 80% credible intervals only.

It may be useful to add some more information to the graphs of nonparametric effects to distinguish more obviously between different covariates. Ways to do so are the specification of a title or the specification of axis labels. Both possibilities are supported by *BayesX* as demonstrated in the following examples (compare Figure 9 for the resulting plots):

```
> b.plotnonp 1, title="Mother body mass index"
> b.plotnonp 1, xlab="bmi" ylab="f_bmi" title="Mother body mass index"
```

By default, *BayesX* displays x- and y-axis with five equidistant ticks according to the range of the data that is to be visualised. These defaults may be overwritten using the options `xlimbottom`, `xlimtop` and `xstep` for the x-axis and `ylimbottom`, `ylimtop` and `ystep` for the y-axis, respectively. The usage of these options is more or less self-explanatory and is demonstrated in the following commands which lead to the graph shown in Figure 10.

```
> b.plotnonp 1, xlab="bmi" ylab="f_bmi" title="Mother body mass index"
  ylimbottom=-0.8 ylimtop=0.6 ystep=0.2 xlimbottom=12 xlimtop=40
```




Figure 9: Specification of title and axis labels.

Figure 10 also includes a graph for the effect of the age of the child that is customised in the same way as for the effect of *bmi*.

```
> b.plotnonp 3, xlab="age" ylab="f_age" title="Age of the child in months"
  ylimbottom=-0.3 ystep=0.3 xlimbottom=0 xlimtop=60 xstep=10
```




Figure 10: Re-defining x- and y-axis.

Now we turn to the options for method **drawmap**. By default, **drawmap** uses grey scales to represent different values of the posterior mean. Using the option **color** forces *BayesX* to use different colours instead. Here the default would be to represent higher values through green colours and smaller values through red colours. Specifying **swapcolors** switches this definition. Therefore the following command

```
> b.drawmap 5, color swapcolors
```

leads to the graph shown in Figure 11 with higher values being represented through red colours and smaller values through green colours. An alternative color scheme can be requested by adding option **hcl**.

Similar options as for the visualisation of nonparametric effects exist for method **drawmap**. For example, a title may be included by specifying the option **title**

```
> b.drawmap 5, color swapcolors title="Structured spatial effect"
```

or the range of values to be displayed may be defined using the options **lowerlimit** and **upperlimit**:




Figure 11: Posterior mean of the structured spatial effect in colour.

```
> b.drawmap 5, color swapcolors title="Structured spatial effect" lowerlimit=-0.3
upperlimit=0.3
```

The graph produced by the second command is shown in Figure 12.

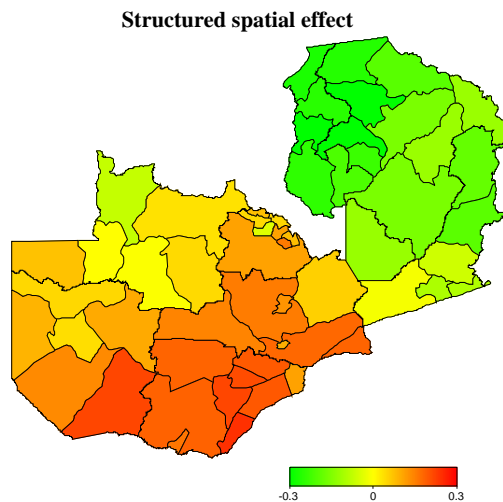


Figure 12: Specifying a title and the range of the plot for spatial effects.

9 Autocorrelation functions and sampling paths

Bayesreg objects provide some post estimation commands to get sampled parameters or to plot autocorrelation functions of sampled parameters. For example

```
> b.plotautocor, maxlag=250
```

computes and displays the autocorrelation functions for all estimated parameters with `maxlag` specifying the maximum lag number (Figure 13 shows a small part of the resulting graph).

If the number of parameters is large, this may be computationally expensive, so *BayesX* provides a second possibility to compute autocorrelation functions. Adding the option `mean` to the `plotautocor` command as in




Figure 13: Autocorrelation function for the scale parameter and the intercept.

```
> b.plotautocor, mean
```

leads to the computation of only the minimum, mean and maximum autocorrelation functions. The result for the scale parameter is shown in Figure 14.

Note, that executing the `plotautocor` command also stores the computed autocorrelation functions in a file named `autocor.raw` in the output directory of the *bayesreg object*. The content of this file can then be visualised in your favorite statistical software package.

To save memory, the sampling paths of the estimated parameters are only stored temporarily by default and will be destroyed, when the corresponding *bayesreg object* is deleted. If we want to store the sampling paths permanently, we have to execute the `getsample` command

```
> b.getsample
```

which stores the sampled parameters in ASCII files in the output directory. To avoid too large files, the samples are typically partitioned into several files. Executing the `getsample` command also produces *ps files* of the sampling paths in the output directory (compare Figure 15 for the content of one of these files).

10 Sensitivity analysis

In some situations, the estimation results of a full Bayesian semiparametric regression model depend on the choice of hyperparameters, e.g. the parameters a and b defining the inverse gamma prior of the variances of nonparametric and spatial effects. Hence, it is often recommended to check how sensitive the results are with respect to changes in the hyperparameters. In the following we will re-estimate the model from section 6 with different choices for the hyperparameters a and b for each effect in the model. The standard choices for a and b are $a = b = 0.001$. As a first trial we choose a smaller value for a and b :

```
> b.regress hazstd = rcw + edu1 + edu2 + tpr + sex
+ bmi(psplinerw2,a=0.00001,b=0.00001) + agc(psplinerw2,a=0.00001,b=0.00001)
+ district(spatial,map=m,a=0.00001,b=0.00001)
+ district(random,a=0.00001,b=0.00001), family=gaussian iterations=12000
burnin=2000 step=10 predict using d
```

Figure 16 shows the results for the nonparametric effects with this choice of hyperparameters. Obviously, the estimated functions are somewhat smoother but they do not differ that much from the estimates with the standard choices.




Figure 14: Minimum, mean and maximum autocorrelation function for the scale parameter.

Now we try two further choices for the hyperparameters, each with $a = 1$ and b small. We estimate models with $b = 0.005$ and $b = 0.00005$:

```
> b.regress hazstd = rcw + edu1 + edu2 + tpr + sex + bmi(psplinerw2,a=1,b=0.005)
+ agc(psplinerw2,a=1,b=0.005) + district(spatial,map=m,a=1,b=0.005)
+ district(random,a=1,b=0.005), family=gaussian iterations=12000 burnin=2000
step=10 predict using d

> b.regress hazstd = rcw + edu1 + edu2 + tpr + sex + bmi(psplinerw2,a=1,b=0.00005)
+ agc(psplinerw2,a=1,b=0.00005) + district(spatial,map=m,a=1,b=0.00005)
+ district(random,a=1,b=0.00005), family=gaussian iterations=12000 burnin=2000
step=10 predict using d
```

Figure 17 and 18 contain the results for the nonparametric effects for the two choices of hyperparameters.

References

- Brezger, A. & Lang, S., 2006: Generalized structured additive regression based on Bayesian P-splines. *Computational Statistics and Data Analysis*, **50**, 967–991.
- Fahrmeir, L., Kneib, T. & Lang, S., 2004: Penalized structured additive regression for space-time data: A Bayesian perspective, *Statistica Sinica*, **14**, 731–761 .




Figure 15: Sampling path of the intercept.




Figure 16: Results for the nonparametric effects with hyperparameters $a = b = 0.00001$ for nonparametric and spatial effects.

- Fahrmeir, L. & Lang, S., 2001a: Bayesian Inference for Generalized Additive Mixed Models Based on Markov Random Field Priors. *Journal of the Royal Statistical Society C*, **50**, 201–220.
- Fahrmeir, L. & Lang, S., 2001: Bayesian Semiparametric Regression Analysis of Multicategorical Time-Space Data. *Annals of the Institute of Statistical Mathematics*, **53**, 10–30
- Fahrmeir, L. & Osuna, L. (2003), Structured additive regression for overdispersed and zero-inflated count data. *Applied Stochastic Models in Business and Industry*, **22**, 351–369.
- George, A. & Liu, J.W. 1981: *Computer Solution of Large Sparse Positive Definite Systems*, Prentice-Hall.
- Hennerfeind, A., Brezger, A. & Fahrmeir, L., 2003: Geoadditive survival models. *Journal of the American Statistical Association*, **101**, 1065–1075
- Kandala, N. B., Lang, S., Klasen, S. & Fahrmeir, L. (2001): Semiparametric Analysis of the Socio-Demographic and Spatial Determinants of Undernutrition in Two African Countries. *Research in Official Statistics*, **1**, 81–100.
- Lang, S. & Brezger, A., 2004: Bayesian P-splines. *Journal of Computational and Graphical Statistics*, **13**, 183–212.
- Spiegelhalter, D.J., Best, N.G., Carlin, B.P. & van der Linde, A. (2002): Bayesian measures of model complexity and fit. *Journal of the Royal Statistical Society B*, **65**, 583–639




Figure 17: Results for the nonparametric effects with hyperparameters $a = 1$ and $b = 0.005$ for nonparametric and spatial effects.




Figure 18: Results for the nonparametric effects with hyperparameters $a = 1$ and $b = 0.00005$ for nonparametric and spatial effects.